



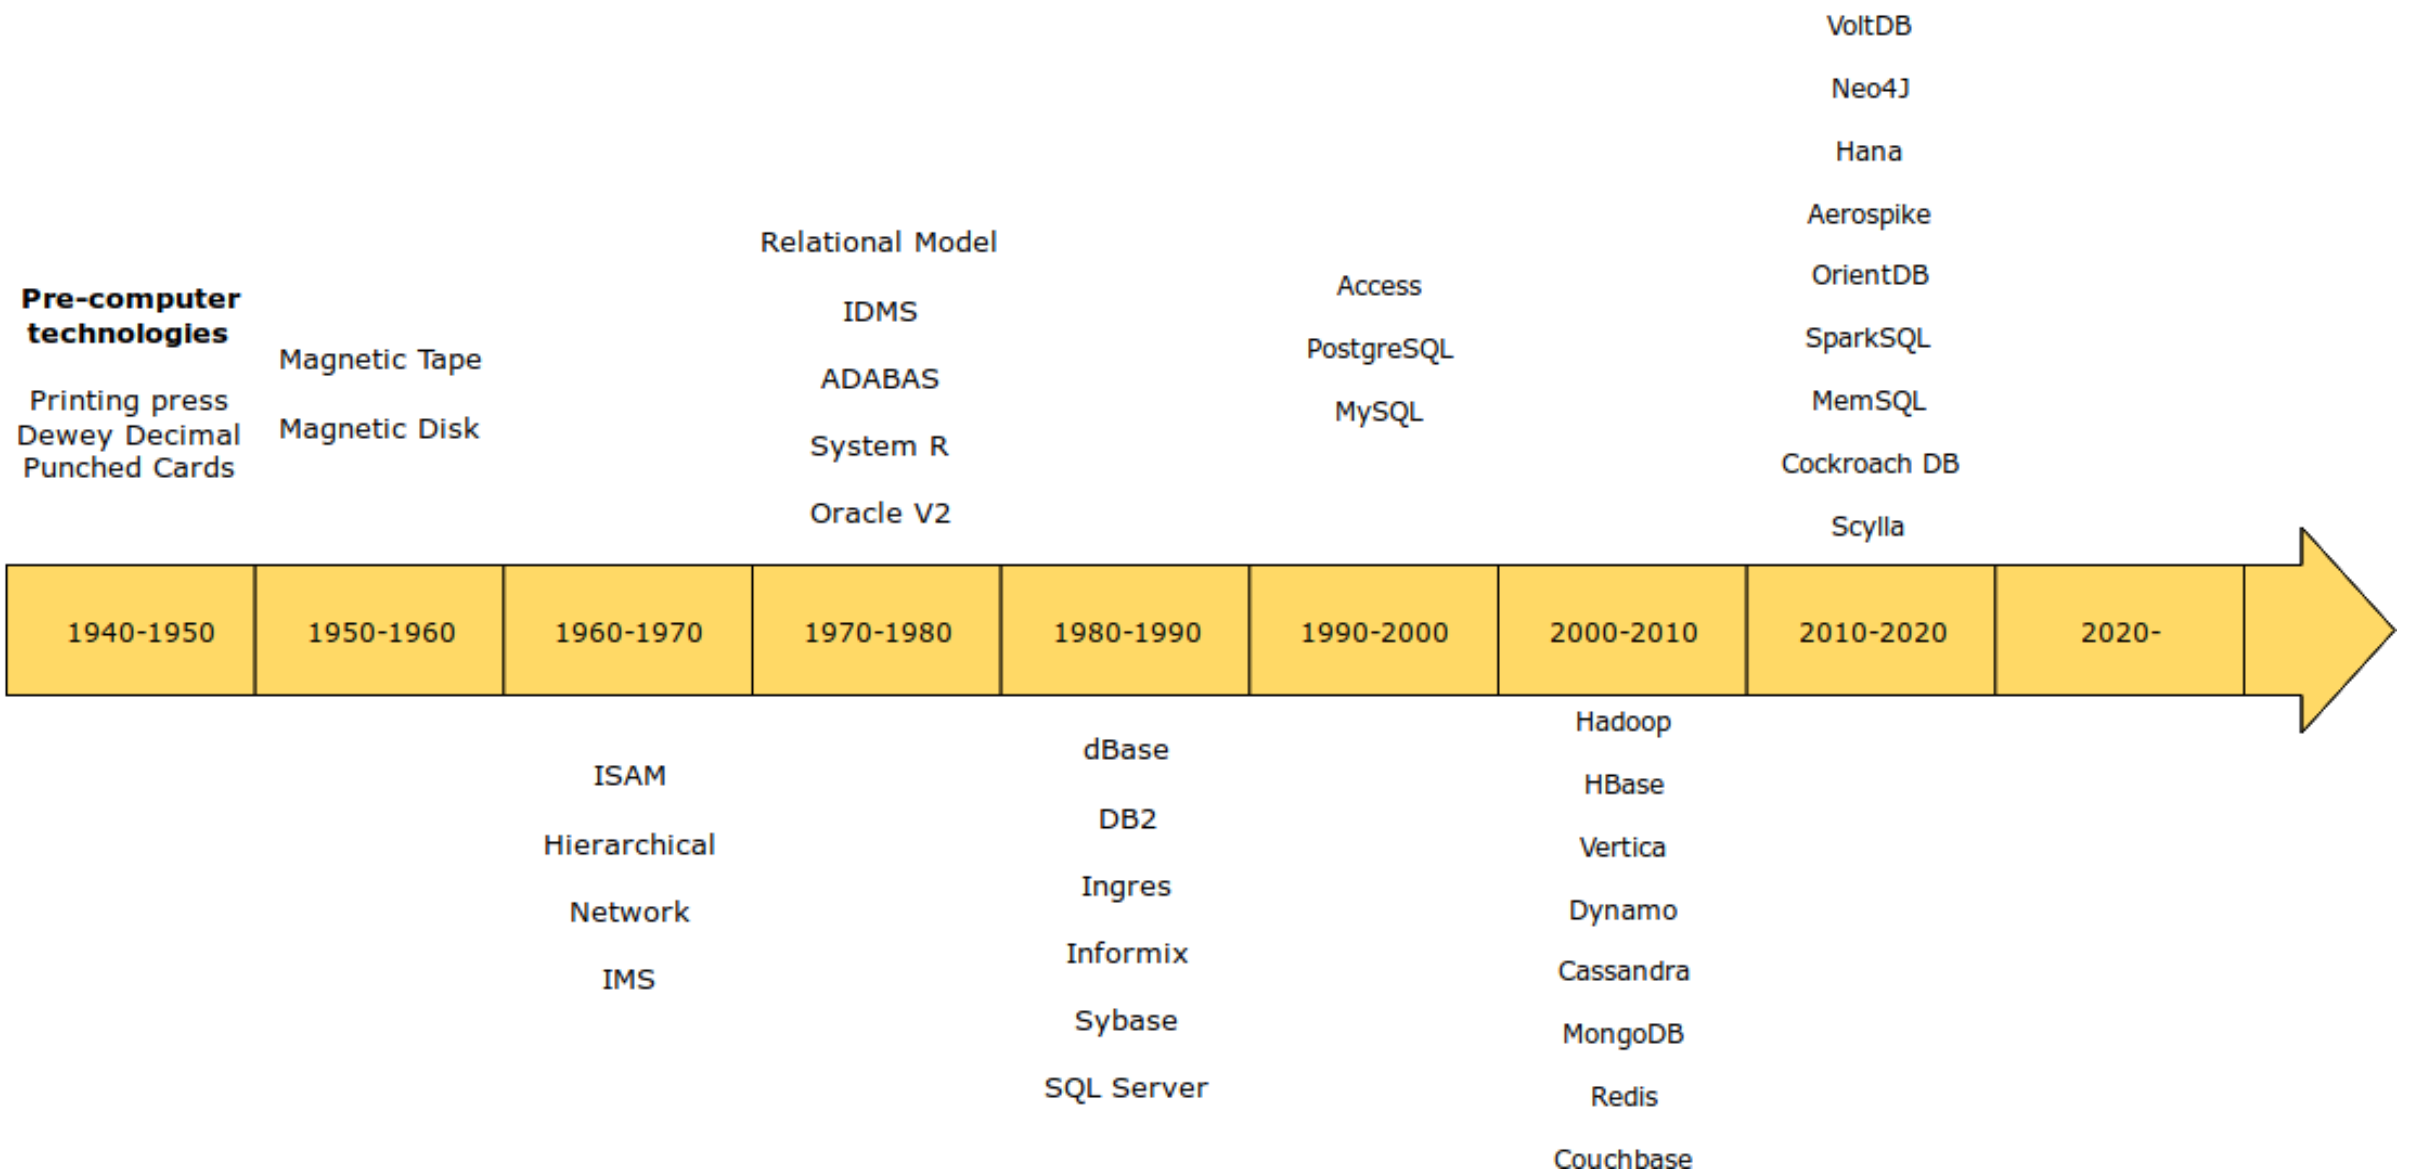
AWS Quick Start

AWS Purpose-Built Database Strategy: The Right Tool for The Right Job

Blair Layton, Business Development Manager, AWS, ASEAN




The Database Market is Changing



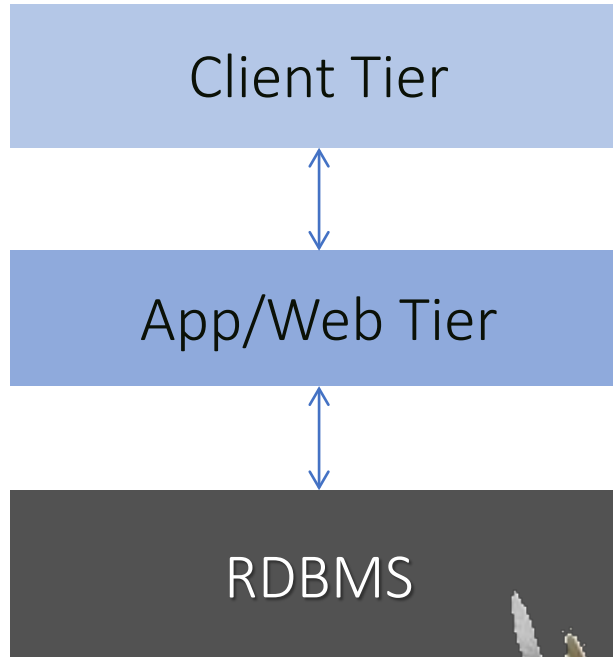
Traditional Database Architecture

*one database for
all workloads*



Traditional Database Architecture

- key-value access
- complex queries
- transactions
- analytics





Data categories and common use cases



Relational

Referential integrity, ACID transactions, schema-on-write

Lift and shift, EMR, CRM, finance



Key-value

Low-latency, key lookups with high throughput and fast ingestion of data

Real-time bidding, shopping cart, social



Document

Indexing and storing documents with support for query on any attribute

Content management, personalization, mobile



In-memory

Microseconds latency, key-based queries, and specialized data structures

Leaderboards, real-time analytics, caching



Graph

Creating and navigating data relations easily and quickly

Fraud detection, social networking, recommendation engine



Search

Indexing and searching semi-structured logs and data

Product catalog, help and FAQs, full text



Time-series

Collect, store, and process data sequenced by time

IoT applications, event tracking




Ledger

Complete, immutable, and verifiable history of all changes to application data

Systems of record, supply chain, health care, registrations, financial

Application Architecture is Changing

Microservices at Amazon



Service-Oriented Architecture (SOA)


Single-purpose

Connect only through APIs



Connect over HTTPS

“Microservices”



Monolithic vs. SOA vs. Microservices



Monolithic
Single Unit



SOA
Coarse-grained



Microservices
Fine-grained



Microservices...

Eliminates any long-term commitment to a technology stack

Polyglot **ecosystem**

Polyglot **persistence**

- Decompose Databases
- **Database per microservice pattern**

Allows easy use of Canary and **Blue-Green** deployments

Purpose-Built Databases



Our portfolio

Broad and deep portfolio, purpose-built for builders

Business Intelligence & Machine Learning



QuickSight



SageMaker

Analytics



Redshift
Data warehousing



EMR
Hadoop + Spark



Athena
Interactive analytics



Elasticsearch Service
Operational Analytics



Kinesis Data Analytics Real time

Databases



Aurora
MySQL, PostgreSQL



DynamoDB
Key value, Document



QLDB NEW
Ledger Database



RDS
MySQL, PostgreSQL, MariaDB,
Oracle, SQL Server



ElastiCache
Redis, Memcached



Timestream NEW
Time Series



RDS on VMware NEW



Neptune
Graph



DocumentDB NEW
Document

Blockchain



Managed Blockchain NEW



Blockchain Templates

Data Lake



S3/Glacier



Lake Formation NEW
Data Lakes



Glue
ETL & Data Catalog

Data Movement

Database Migration Service | Snowball | Snowmobile | Kinesis Data Firehose | Kinesis Data Streams





AWS: Purpose-built databases




Relational



Key-value



Document



In-memory



Graph



Search



Time-series



Ledger



Amazon RDS



Amazon
DynamoDB



Amazon
DocumentDB



Amazon
ElastiCache



Amazon
Neptune



Amazon
Elasticsearch
Service



Amazon
Timestream



Amazon
Quantum
Ledger
Database

Aurora Community Commercial



ORACLE



Microsoft
SQL Server



Relational Databases



Amazon RDS

Managed relational database service with a choice of popular database engines

Amazon
Aurora




Microsoft SQL Server

ORACLE




Easy to administer

No need to provision infrastructure, install, and maintain DB software



Available & durable

Automatic Multi-AZ data replication; automated backup, snapshots, and failover



Highly scalable

Scale DB compute and storage with a few clicks; minimal downtime for your application



Fast & secure

SSD storage and guaranteed provisioned I/O; data encryption at rest and in transit



Amazon Aurora


MySQL and PostgreSQL compatible relational database built for the cloud

Performance and availability of commercial-grade databases at 1/10th the cost




Performance & scalability

5x throughput of standard MySQL and 3x of standard PostgreSQL; scale-out up to 15 read replicas




Availability & durability

Fault-tolerant, self-healing storage; six copies of data across three AZs; continuous backup to S3



Highly secure

Network isolation, encryption at rest/transit



Fully managed

Managed by RDS: no hardware provisioning, software patching, setup, configuration, or backups

Aurora: Fastest Growing Service in AWS History

NASDAQ

SAMSUNG

FINRA

bmc

URBAN AIRSHIP

EA
ELECTRONIC ARTS™

hulu

nielsen

Pearson

ancestry

RavenPack

zynga

ZUMBA
FITNESS

Blackboard

THOMSON REUTERS

SAFE SOFTWARE™

INRIX

Expedia®

DOW JONES

INTERCOM

Alfresco™

THE HONEST CO.
THE HONEST CO.

FICO®

jenny
CRAIG

AstraZeneca

CBS Interactive

workfront

gumi
One Step Beyond.

Experian™

delight
media&camera works

Zillow®

Lookout

NTT
docomo

PeopleAdmin

rackspace®

FUNNY
& DIE

SysAid

Bristol-Myers Squibb

毎日新聞

CyberAgent®

SRA OSS

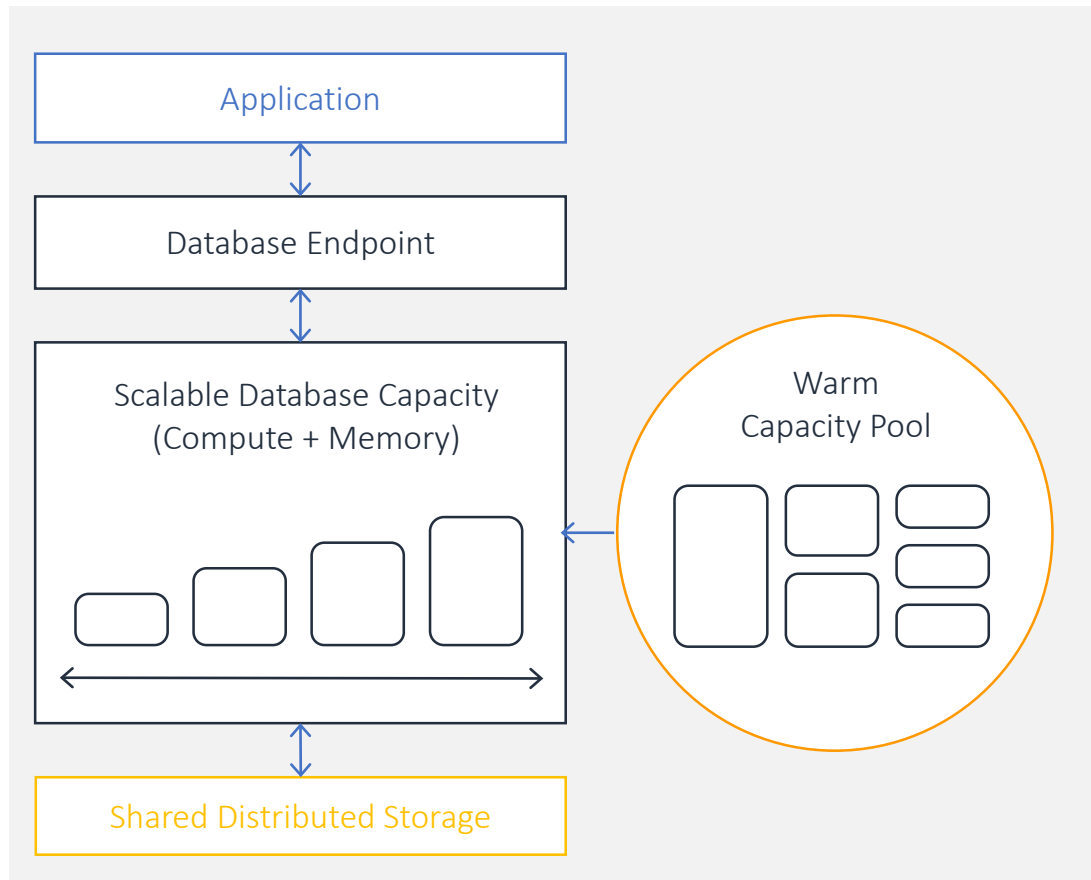
TalentBin
SOURCE THE WEB.

aws



Aurora Serverless

On-demand, auto-scaling database for applications with variable workloads



Starts up on demand, shuts down when not in use

Automatically scales with no instances to manage

Pay per second for the database capacity you use

Non-Relational Databases

Amazon DynamoDB

We needed to adapt to power Amazon.com

Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall and Werner Vogels

Amazon.com

ABSTRACT

Reliability at massive scale is one of the biggest challenges we face at Amazon.com, one of the largest e-commerce operations in the world; even the slightest outage has significant financial consequences and impacts customer trust. The Amazon.com platform, which provides services for many web sites worldwide, is implemented on top of an infrastructure of tens of thousands of servers and network components located in many data centers around the world. At this scale, small and large components fail continuously and the way persistent state is managed in the face of these failures drives the reliability and availability of the software systems.

This paper presents the design and implementation of Dynamo, a highly available key-value storage system that some of Amazon's core services use to provide an "always-on" experience. To achieve this level of availability, Dynamo sacrifices consistency under certain failure scenarios. It makes extensive use of object versioning and application-assisted conflict resolution in a manner that provides a novel interface for developers to use.

One of the lessons our organization has learned from operating Amazon's platform is that the reliability and scalability of a system is dependent on how its application state is managed. Amazon.com is a highly decentralized, loosely coupled, service oriented architecture consisting of hundreds of services. In this environment there is a particular need for storage technologies that are always available. For example, customers should be able to view and add items to their shopping cart even if disks are failing, network routes are flapping, or data centers are being destroyed by tornados. Therefore, the service responsible for managing shopping carts requires that it can always write to and read from its data store, and that its data needs to be available across multiple data centers.

Dealing with failures in an infrastructure comprised of millions of components is our standard mode of operation; there are always a small but significant number of server and network components that are failing at any given time. As such Amazon's software systems need to be constructed in a manner that treats failure handling as the normal case without impacting availability or performance.

Needed to power Amazon.com

Required massive scalability and reliability

DynamoDB designed to meet this need



Many Applications Require Millisecond Latency at Any Scale

Example: Amazon Prime Day 2017



Biggest shopping event in Amazon history

Thousands of Amazon teams using DynamoDB

3.34 trillion requests

Peaked at **12.9 million** requests per second




Amazon DynamoDB

Fully-managed nonrelational database for any scale




Fully managed

- Maintenance-free
- Serverless
- Auto scaling
- Backup and restore
- Global tables



High performance


- Fast, consistent performance
- Virtually unlimited throughput
- Virtually unlimited storage



Secure

- Encryption at rest and transit
- Fine-grained access control
- PCI, HIPAA, FIPS140-2 eligible

DynamoDB Global Tables



Build high-performance, globally distributed applications

Low latency reads and writes to locally available tables

Multi-region redundancy and resiliency

Easy to set up and no application rewrites required



Amazon ElastiCache

Fully-managed, Redis or Memcached compatible, low-latency, in-memory data store




Extreme performance

In-memory data store and cache for sub-millisecond response times



Fully managed

AWS manages all hardware and software setup, configuration, monitoring




Easily scalable

Read scaling with replicas
Write and memory scaling with sharding
Non-disruptive scaling

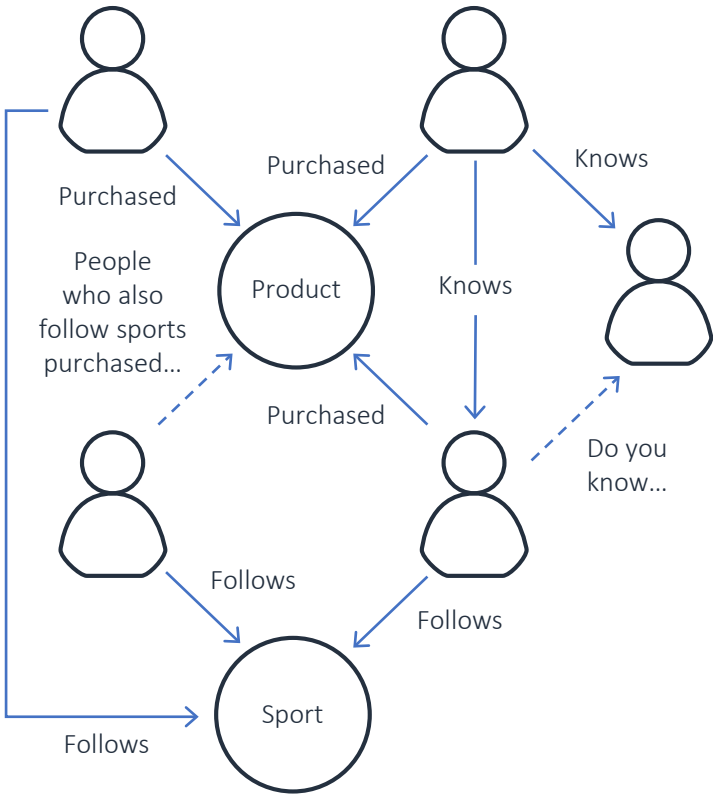


Graph Use Cases

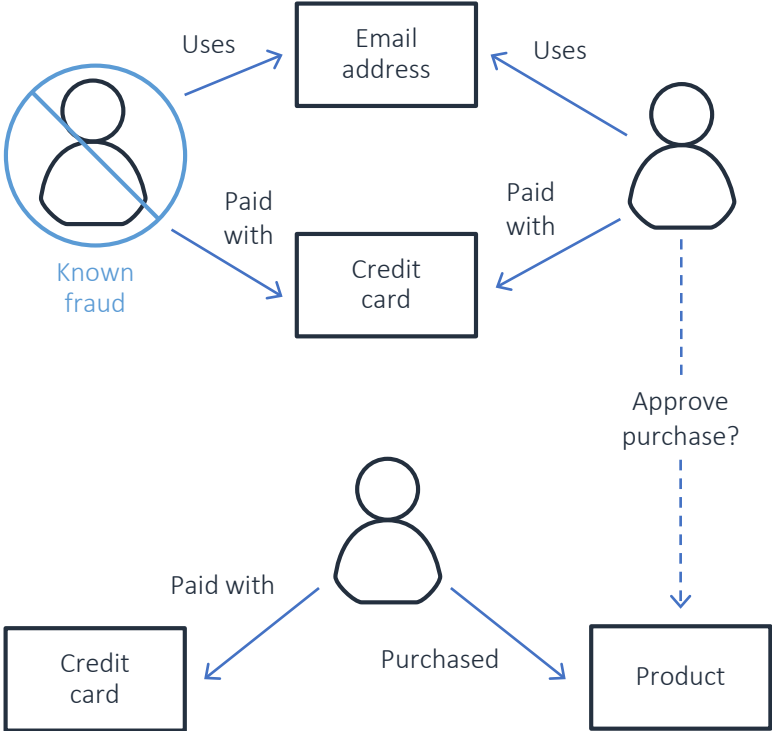
Social news feed



Recommendations



Retail fraud detection





Challenges Building Apps with Highly-Connected Data


Relational databases



Unnatural for querying graph




Inefficient graph processing



Rigid schema inflexible for changing graphs

Existing graph databases




Difficult to scale



Difficult to maintain high availability



Too expensive




Limited support for open standards



Amazon Neptune

Fully managed graph database




Open

Supports Apache TinkerPop and W3C RDF graph models



Fast & scalable

Store billions of relationships; query with millisecond latency



Reliable

Six replicas of your data across three AZs with full backup and restore



Easy

Build powerful queries easily with Gremlin and SPARQL



Amazon DocumentDB

Fast, scalable, and fully managed MongoDB-compatible database service

Fast



Millions of requests per second with millisecond latency; twice the throughput of MongoDB

Scalable



Separation of compute and storage enables both layers to scale independently; scale out to 15 read replicas in minutes

Fully managed



Managed by AWS: no hardware provisioning; auto patching, quick setup, secure, and automatic backups

MongoDB compatible



Compatible with MongoDB 3.6; use the same SDKs, tools, and applications with Amazon DocumentDB

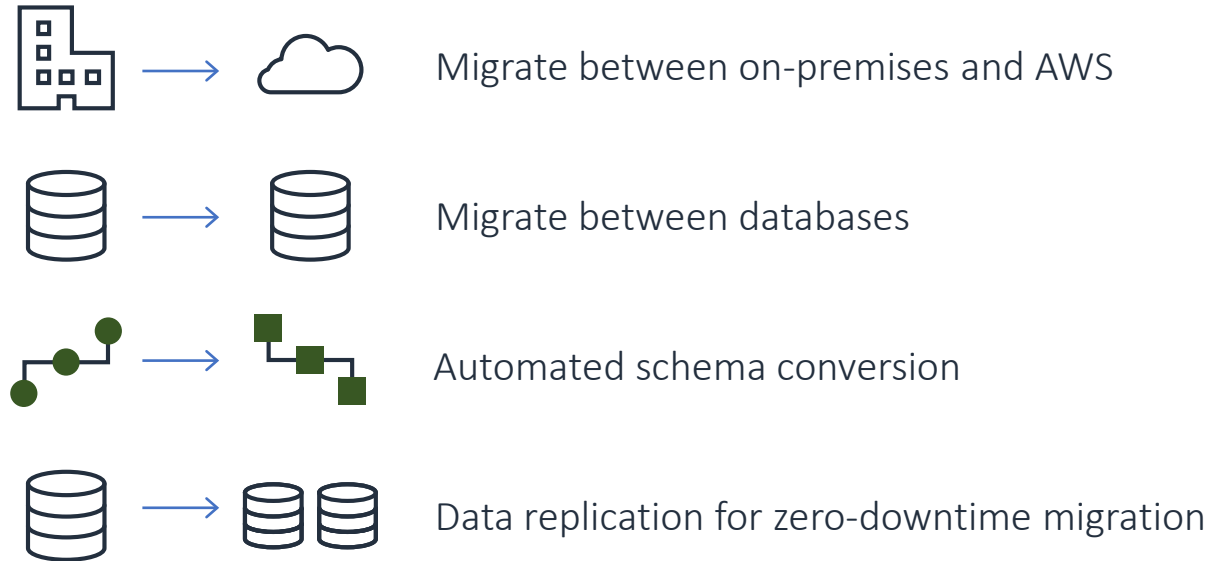
Database Migration Service



AWS Database Migration Service

Migrating Databases to AWS

100,000+
Databases migrated



100,000+ Databases Migrated with DMS





Thank You for Attending AWS Quick Start

We hope you found it interesting! A kind reminder to **complete the survey**.
Let us know what you thought of today's event and how we can improve the event experience for you in the future.



aws-apac-marketing@amazon.com



twitter.com/AWSCloud



facebook.com/AmazonWebServices



youtube.com/user/AmazonWebServices



slideshare.net/AmazonWebServices



twitch.tv/aws