

# **AWS INNOVATE**

---

## ONLINE CONFERENCE 2019

---

### MACHINE LEARNING & AI EDITION

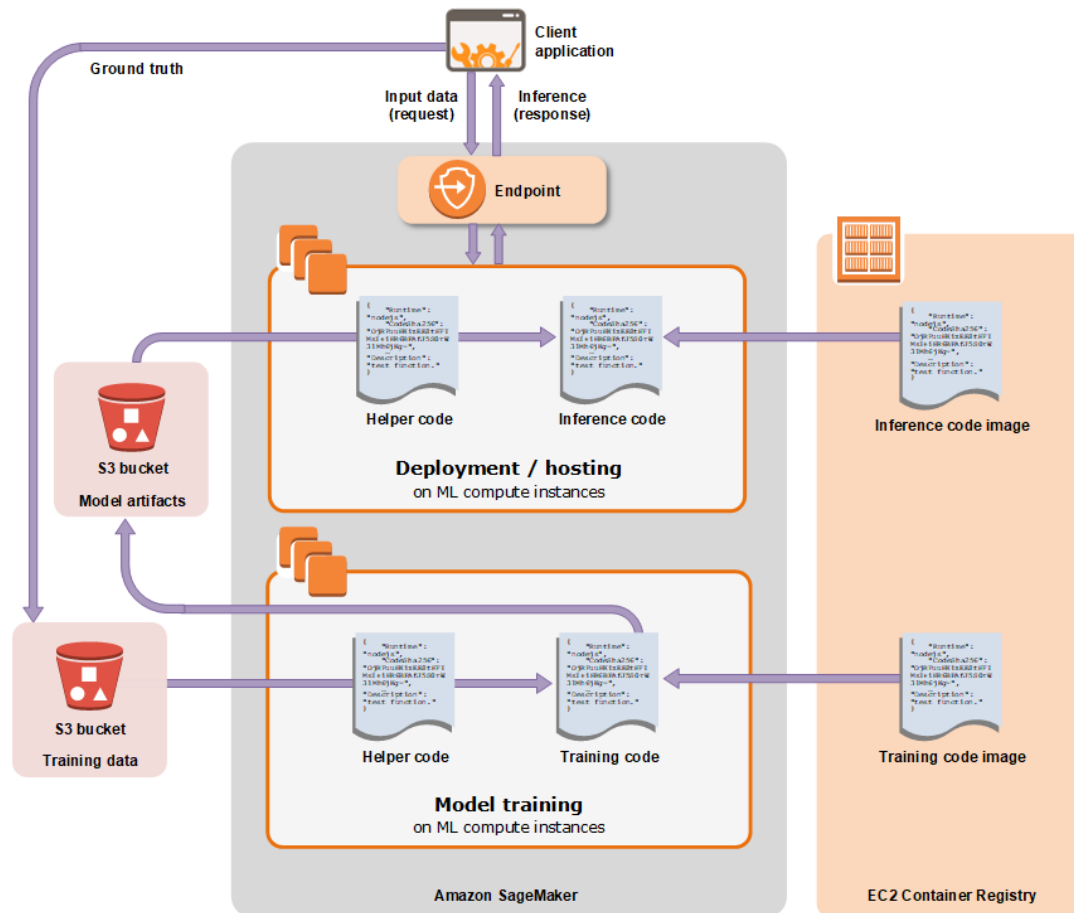
## **Opensource ML Frameworks on Amazon SageMaker Featuring TensorFlow, PyTorch, MXnet, Keras and Gluon (Level 300)**

Aparna Elangovan, Solutions Architect, AWS ANZ

# Deep Learning on Amazon SageMaker

- Built-in support for TensorFlow, PyTorch, Chainer, MXNet
- Fully customizable Bring Your Own (BYO) container option
- Distributed GPU training
- Code Portability
- Experiment Tracking
- One-click training and One-click deployment

# Amazon SageMaker – Training phase



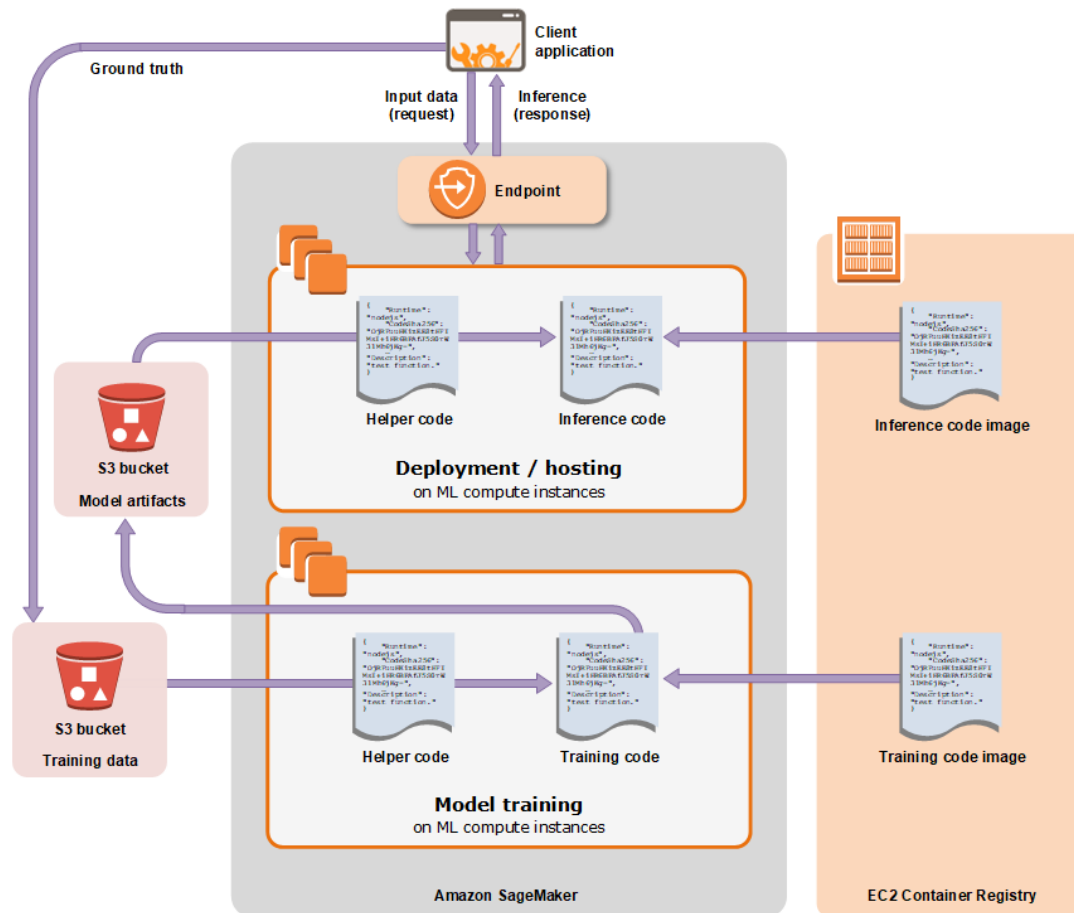
## You provide

1. Your source code in S3 or local file system
2. Your data usually in S3
3. A base docker image, SageMaker Built-in or BYO
4. The EC2 instance type and count

## Amazon SageMaker

- ✓ Automatically provisions and tears down Amazon EC2 instances for your training job
- ✓ Launches docker image in the EC2 instance type required for training & inference
- ✓ Downloads code & data from S3 to container host
- ✓ Executes training job
- ✓ Copies model from container host to S3
- ✓ Copies results from container host to S3
- ✓ Logs / print results captured in Amazon CloudWatch
- ✓ Captures training infrastructure & performance metrics in CloudWatch

# Amazon SageMaker – Inference phase



## You provide

1. The trained model
2. The docker Image
3. The EC2 instance type and initial count

## SageMaker

- ✓ Launches docker image in the EC2 instance type
- ✓ Downloads model from S3 to container host
- ✓ Auto scales instances based on scaling policy
- ✓ Automatically captures metrics such as latency, #invocations
- ✓ Logs / print results captured in CloudWatch

# Hyperparameters and Environment Variables

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the training directory", required=True)
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts such as training data',
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input file wrt to the val directory", required=True)
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifacts such as validation data',
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "."))


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "result_data")
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by SageMaker", default=None)

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the snapshot to..",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10, type=int)
parser.add_argument("--batch-size", help="The mini batch size", default=30, type=int)
```

main\_train.py



# Hyperparameters and Environment Variables

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
```

```
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
```

```
abalone_estimator = TensorFlow(entry_point='main_train.py',
                               source_dir='./source',
                               role=role,
                               py_version="py3",
                               framework_version = "1.11.0",
                               hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                               model_dir = s3_model_path,
                               metric_definitions = [{"Name": "mean_squared_error",
                                                       "Regex": "## validation_metric_mse ##: (\d*[\.]\d*)"},
                                                     {"Name": "mean_absolute_error",
                                                       "Regex": "## validation_metric_mae ##: (\d*[\.]\d*)"},
                                                     {"Name": "mean_absolute_percentage_error",
                                                       "Regex": "## validation_metric_mape ##: (\d*[\.]\d*)"}],
                               train_instance_count=1,
                               train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit( {'train': s3_input_prefix,
                        'validation': s3_input_prefix},
                        job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the tra
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input fi
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifact
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save result
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "resu
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=
parser.add_argument("--batch-size", help="The mini batch size", defau
```

main\_train.py





# Hyperparameters and Environment Variables

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the tr
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input fi
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifac
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save resul
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "res
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=
parser.add_argument("--batch-size", help="The mini batch size", defau
```

main\_train.py



## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
```

```
s3_model_path = "s3://{}/models".format(sagemaker_session.d
```

```
abalone_estimator = TensorFlow(entry_point='main_train.py',
                               source_dir='./source',
                               role=role,
                               py_version="py3",
                               framework_version = "1.11.0",
                               hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                               model_dir = s3_model_path,
                               metric_definitions = [{"Name": "mean_squared_error",
                                                       "Regex": "## validation_metric_mse ##: (\d*[\.]?\d*)"},
                                                     {"Name": "mean_absolute_error",
                                                       "Regex": "## validation_metric_mae ##: (\d*[\.]?\d*)"},
                                                     {"Name": "mean_absolute_percentage_error",
                                                       "Regex": "## validation_metric_mape ##: (\d*[\.]?\d*)"}
                                                    ],
                               train_instance_count=1,
                               train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit( {'train': s3_input_prefix,
                        'validation':s3_input_prefix},
                        job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

## 1) Specify source code

The entry point file and source code dir

# Hyperparameters and Environment Variables

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
```

```
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
```

```
abalone_estimator = TensorFlow(entry_point='main_train.py',
                               source_dir='./source',
                               role=role,
                               py_version="py3",
                               framework_version = "1.11.0",
                               hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                               model_dir = s3_model_path,
                               metric_definitions = [{"Name": "mean_validation_rmse",
                                                       "Regex": "## validation rmse"},
                                                      {"Name": "mean_validation_rmse",
                                                       "Regex": "## validation rmse"},
                                                      {"Name": "mean_validation_rmse",
                                                       "Regex": "## validation rmse"}],
                               train_instance_count=1,
                               train_instance_type='ml.c4.xlarge')
```

## 2) Map S3 prefix to local download directory

The key name e.g. 'train' matches environment variable **SM\_CHANNEL\_TRAIN** suffix **TRAIN**, for train data directory

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="abalone_age_py3_{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the train data")
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts',
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))
# val dir files
parser.add_argument('--validationdata', help="The validation input file")
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifacts',
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "."))


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "results"))

parser.add_argument("--model_dir", help="Do not use this.. required by sagemaker")

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the model",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

main\_train.py





# Hyperparameters and Environment Variables

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the tra
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input fil
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifact
                    default=os.environ.get('SM_CHANNEL_VALIDATION', ".


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save result
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "resu
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=
parser.add_argument("--batch-size", help="The mini batch size", default=
```

main\_train.py



Submit your training job

```
from sagemaker.tensorflow import TensorFlowTrainer
from time import gmtime, strftime
```

```
s3_model_path = "s3://{}/models".format(bucket, job_name)
```

```
abalone_estimator = TensorFlowTrainer(
    source_dir="./source",
    role=role,
    py_version="py3",
    framework_version="1.11.0",
```

```
hyperparameters={
    'traindata': 'abalone_train.csv',
    'validationdata': 'abalone_test.csv',
    'epochs': 10,
    'batch-size': 32,
}
```

```
model_dir=s3_model_path,
metric_definitions=[{"Name": "mean_squared_error",
                    "Regex": "## validation_metric_mse ##: (\d*[\.]\d*)"},
                    {"Name": "mean_absolute_error",
                    "Regex": "## validation_metric_mae ##: (\d*[\.]\d*)"},
                    {"Name": "mean_absolute_percentage_error",
                    "Regex": "## validation_metric_mape ##: (\d*[\.]\d*)"}],
train_instance_count=1,
```

```
train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({
    'train': s3_input_prefix,
    'validation': s3_input_prefix,
```

```
job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

## 3) Pass hyper parameters

The hyperparameter dict key name e.g. "traindata", "epochs" matches the argument name "--traindata", "--epochs"

# Hyperparameters and Environment Variables

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
```

```
s3_model_path =
```

```
abalone_estimator =
```

### 4) Save artifacts to output

Save model to output to dir pointed by environment variable SM\_MODEL\_DIR. Artifacts placed here are automatically uploaded to S3 and available during inference

```
validationdata = abalone_test.csv,
'epochs': 10,
```

```
batch-size = 32,
```

```
model_dir = s3_model_path,
```

```
metric_definitions = [{"Name": "mean_squared_error",
```

```
"Regex": "## validation_metric_mse ##: (\d*[\.]\d*)"}],
```

```
{ "Name": "mean_absolute_error",
```

```
"Regex": "## validation_metric_mae ##: (\d*[\.]\d*)"}],
```

```
{ "Name": "mean_absolute_percentage_error",
```

```
"Regex": "## validation_metric_mape ##: (\d*[\.]\d*)"}],
```

```
],
```

```
train_instance_count=1,
```

```
train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit( {'train': s3_input_prefix,
'validation':s3_input_prefix},
```

```
job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime()))))
```

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the tr
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input fi
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifac
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save resul
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "res
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=
parser.add_argument("--batch-size", help="The mini batch size", defau
```

main\_train.py



# Hyperparameters and Environment Variables

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
```

```
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
```

```
abalone_estimator = TensorFlow(entry_point='main_train.py',
                               source_dir='./source',
                               role=role,
                               py_version="py3",
                               framework_version = "1.11.0",
                               hyperparameters={'traindata': 'abalone',
                                                  'validationdata': 'abalone',
                                                  'epochs': 10,
                                                  'batch-size': 32},
                               model_dir = s3_model_path,
                               metric_definitions = [{"Name": "validation_metric_mse",
                                                       "Regex": "## validation_metric_mse ##: (\d*[\.]? \d*)"},
                                                      {"Name": "mean_absolute_error",
                                                       "Regex": "## validation_metric_mae ##: (\d*[\.]? \d*)"},
                                                      {"Name": "mean_absolute_percentage_error",
                                                       "Regex": "## validation_metric_mape ##: (\d*[\.]? \d*)"}],
                               train_instance_count=1,
                               train_instance_type='ml.c4.xlarge')
```

5) Specify the instance type for your training. Increase the instance count if your code supports distributed training

```
abalone_estimator.fit( {'train': s3_input_prefix,
                        'validation': s3_input_prefix},
                        job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the tr
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input fi
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifac
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "


# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save resul
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "res
                    )

parser.add_argument("--model_dir", help="Do not use this.. required b

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=
parser.add_argument("--batch-size", help="The mini batch size", defau
```

main\_train.py



# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(training_dir, training_filename))
val_x, val_y = input_transformer_load(os.path.join(val_dir, val_filename))
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, validation_data=(val_x, val_y))
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone_age_predictor.h5")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it can be extracted using a regular express
# This allows SageMaker to report this metrics and allows hyper parameter tuning
# Note: Use a special marker for SageMaker to extract the metrics, say ## Metric ##
for i, m in enumerate(metrics):
    print("## validation_metric_{0} ##: {1}".format(m, scores[1+i]))
```

- source
  - main\_train.py
  - model\_exporter\_keras\_to\_pb.py
  - requirements.txt
  - setup.py
- tests
  - \_\_init\_\_.py
  - test\_train.py

main\_train.py

# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(tr
val_x, val_y = input_transformer_load(os.path.join(val_dir
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size)
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone.keras")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it
# This allows SageMaker to report this metrics and allows
# Note: Use a special marker for SageMaker to extract the
for i, m in enumerate(metrics):
    print("## validation_metric_{} ##: {}".format(m, scores[i]))
```

```
▼ source
  main_train.py
  model_exporter_keras_to_pb.py
  requirements.txt
  setup.py
▼ tests
  __init__.py
  test_train.py
```

main\_train.py

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = TensorFlow(entry_point='main_train.py',
                              source_dir="./source",
                              role=role,
                              py_version="py3",
                              framework_version="1.11.0",
                              hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                              model_dir=s3_model_path,
                              metric_definitions=[{"Name": "mean_squared_error",
                                                    "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_error",
                                                    "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_percentage_error",
                                                    "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                              train_instance_count=1,
                              train_instance_type='ml.c4.xlarge')

abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix,
                      job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime()))})
```



# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(tr
val_x, val_y = input_transformer_load(os.path.join(val_dir
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size)
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone_keras.h5")
model.save(model_keras_path)
```

```
# Step 2: Log your metrics in a special format so that it
# This allows SageMaker to report this metrics and allows
# Note: Use a special marker for SageMaker to extract the
for i, m in enumerate(metrics):
    print("## validation metric {} ##: {}".format(m, score))
```

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = TensorFlow(entry_point=train.py,
                               source_dir="./source",
                               role=role,
                               py_version="py3",
                               framework_version="1.11",
                               hyperparameters={'traindata': train_data_path,
                                                'validationdata': validation_data_path,
                                                'epochs': 10,
                                                'batch-size': 32},
                               model_dir=s3_model_path)
```

## 6) Model metrics regex to trace and visualize your model performance

```
metric_definitions = [{"Name": "mean_squared_error",
    "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
    {"Name": "mean_absolute_error",
    "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
    {"Name": "mean_absolute_percentage_error",
    "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}]
```

```
train_instance_count=1,  
train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="abalone-agg-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

```

▼ source
  main_train.py
  model_exporter_keras_to_pb.py
  requirements.txt
  setup.py
▼ tests
  __init__.py
  test_train.py

```

main\_train.py



# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(tr
val_x, val_y = input_transformer_load(os.path.join(val_dir
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size)
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone.keras")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it
# This allows SageMaker to report this metrics and allows
# Note: Use a special marker for SageMaker to extract the
for i, m in enumerate(metrics):
    print("## validation_metric_{} ##: {}".format(m, scores[i]))
```

```
▼ source
  main_train.py
  model_exporter_keras_to_pb.py
  requirements.txt
  setup.py
▼ tests
  __init__.py
  test_train.py
```

main\_train.py

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(bucket)
abalone_estimator = TensorFlow(entry_point="main_train.py",
                               source_dir="./source",
                               role=role,
```

```
                               py_version="py3",
                               framework_version="1.11.0",
                               hyperparameters={
                                   'traindata': 'abalone_train.csv',
                                   'validationdata': 'abalone_test.csv',
                                   'epochs': 10,
                                   'batch-size': 32},
                               model_dir=s3_model_path,
                               metric_definitions=[{"Name": "mean_squared_error",
                                                       "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                       {"Name": "mean_absolute_error",
                                                       "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                       {"Name": "mean_absolute_percentage_error",
                                                       "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                               train_instance_count=1,
                               train_instance_type='ml.c4.xlarge')
```

7) Specify python version (py3)  
TensorFlow framework version  
(1.11.0)

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(training_dir, training_filename))
val_x, val_y = input_transformer_load(os.path.join(val_dir, val_filename))
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, validation_data=(val_x, val_y))
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone_age_predictor.h5")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it can be extracted using a regular express
# This allows SageMaker to report this metrics and allows hyper parameter tuning
# Note: Use a special marker for SageMaker to extract the metrics, say ## Metric ##
for i, m in enumerate(metrics):
    print("## validation_metric_{} ##: {}".format(m, scores[1+i]))
```

- source
  - main\_train.py
  - model\_exporter\_keras\_to\_pb.py
  - requirements.txt
  - setup.py
- tests
  - \_\_init\_\_.py
  - test\_train.py

main\_train.py

# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(tr
val_x, val_y = input_transformer_load(os.path.join(val_dir
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size)
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone.keras")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it
# This allows SageMaker to report this metrics and allows
# Note: Use a special marker for SageMaker to extract the
for i, m in enumerate(metrics):
    print("## validation_metric_{} ##: {}".format(m, scores[i]))
```

```
▼ source
  main_train.py
  model_exporter_keras_to_pb.py
  requirements.txt
  setup.py
▼ tests
  __init__.py
  test_train.py
```

main\_train.py

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = TensorFlow(entry_point='main_train.py',
                              source_dir="./source",
                              role=role,
                              py_version="py3",
                              framework_version="1.11.0",
                              hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                              model_dir=s3_model_path,
                              metric_definitions=[{"Name": "mean_squared_error",
                                                    "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_error",
                                                    "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_percentage_error",
                                                    "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"},
                                                  ],
                              train_instance_count=1,
                              train_instance_type='ml.c4.xlarge')

abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix,
                      job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime()))})
```

# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(tr
val_x, val_y = input_transformer_load(os.path.join(val_dir
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size)
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone.keras")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it
# This allows SageMaker to report this metrics and allows
# Note: Use a special marker for SageMaker to extract the
for i, m in enumerate(metrics):
    print("## validation_metric_{} ##: {}".format(m, scores[i]))
```

```
▼ source
  main_train.py
  model_exporter_keras_to_pb.py
  requirements.txt
  setup.py
▼ tests
  __init__.py
  test_train.py
```

main\_train.py

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## Submit your training job

```
from sagemaker.tensorflow import TensorFlow
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = TensorFlow(entry_point="main_train.py",
```

```
    source_dir="./source",
    role=role,
    py_version="py3",
    framework_version="1.15",
    hyperparameters={'train': s3_input_prefix,
                     'validation': s3_input_prefix,
                     'epochs': 10,
                     'batch-size': 32},
    model_dir=s3_model_path)
```

```
metric_definitions = [{"Name": "mean_squared_error",
                        "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                       {"Name": "mean_absolute_error",
                        "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                       {"Name": "mean_absolute_percentage_error",
                        "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}]
```

```
train_instance_count=1,
train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix,
                      'job_name': "abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime()))})
```

6) Model metrics regex to trace and visualize your model performance

# Model metrics & Framework Versions

```
# For a mean squared error regression problem
# Using RMSProp optimiser with mean squared error
metrics = ['mse', 'mae', 'mape']
model.compile(optimizer='rmsprop',
              loss='mse', metrics=metrics)

# load train & test data
train_x, train_y = input_transformer_load(os.path.join(tr
val_x, val_y = input_transformer_load(os.path.join(val_dir
# Start training
model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size)
# model evaluate
scores = model.evaluate(val_x, val_y)
# model save in keras default format
model_keras_path = os.path.join(model_snapshotdir, "abalone.keras")
model.save(model_keras_path)

# Step 2: Log your metrics in a special format so that it
# This allows SageMaker to report this metrics and allows
# Note: Use a special marker for SageMaker to extract the
for i, m in enumerate(metrics):
    print("## validation_metric_{} ##: {}".format(m, scores[i]))
```

```
▼ source
  main_train.py
  model_exporter_keras_to_pb.py
  requirements.txt
  setup.py
▼ tests
  __init__.py
  test_train.py
```

main\_train.py

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## Submit your training job

```
from sagemaker.tensorflow import TensorFlowTrainingJob
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(bucket)
abalone_estimator = TensorFlowTrainingJob(
    source_dir="./source",
    role=role,
```

```
    py_version="py3",
    framework_version="1.11.0",
    hyperparameters={
        'traindata': 'abalone_train.csv',
        'validationdata': 'abalone_test.csv',
        'epochs': 10,
        'batch-size': 32},
    model_dir=s3_model_path,
    metric_definitions=[{"Name": "mean_squared_error",
                        "Regex": "## validation_metric_mse ##: (\\d*\\.?)\\d*"}],
    train_instance_count=1,
    train_instance_type='ml.c4.xlarge')
```

7) Specify python version (py3)  
TensorFlow framework version  
(1.11.0)

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix,
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime()))})
```

# Demo – Abalone age prediction

SageMaker with TensorFlow, Keras in Python 3



# PyTorch on SageMaker

# PyTorch training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the training directory", required=True)
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts such as training data',
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input file wrt to the val directory", required=True)
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifacts such as validation data',
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "."))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "result_data")
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by SageMaker", default=None)

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the snapshot to..",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10, type=int)
parser.add_argument("--batch-size", help="The mini batch size", default=30, type=int)
```

# PyTorch training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', ".")

# val dir files
parser.add_argument("--validationdata", help="The validation input",
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', ".")

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', ".")
                    )

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker",
                    default=os.environ.get('SM_MODEL_DIR', None))

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save snapshots",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=100)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.pytorch import PyTorch
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = PyTorch(entry_point='main_train.py',
                             source_dir="./source",
                             role=role,
                             py_version="py3",
                             framework_version="1.0.0",
                             hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                             model_dir=s3_model_path,
                             metric_definitions=[{"Name": "mean_squared_error",
                                                    "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_error",
                                                    "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_percentage_error",
                                                    "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                             train_instance_count=1,
                             train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                       'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# PyTorch training with SageMaker

1) Specify source code  
The entry point file and source code dir

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input",
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', "."))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "."))

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker")

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save snapshots",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=100)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.pytorch import PyTorch
from time import gmtime, strftime
s3_model_path = "s3://[BUCKET]/models" format(sagemaker_session.default_bucket())
abalone_estimator = PyTorch(entry_point='main_train.py',
                             source_dir='./source',
                             role=role,
                             py_version="py3",
                             framework_version="1.0.0",
                             hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                             model_dir=s3_model_path,
                             metric_definitions=[{"Name": "mean_squared_error",
                                                  "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_error",
                                                  "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_percentage_error",
                                                  "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"},
                                                  ],
                             train_instance_count=1,
                             train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# PyTorch training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', ""))

# val dir files
parser.add_argument('--validationdata', help='The validation input',
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', ""))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', ""))

parser.add_argument("--model_dir", help="Do not use this.. requires",
                    # This is where the model needs to be saved to
                    parser.add_argument("--snapshot_dir", help="The directory to save snapshots",
                                        default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.pytorch import PyTorch
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = PyTorch(entry_point='main_train.py',
                             source_dir="./source",
                             role=role,
                             py_version="py3",
                             framework_version="1.0.0",
                             hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                             model_dir=s3_model_path,
                             metric_definitions=[{"Name": "train_loss",
                                                    "Regex": ".*loss.*",
                                                    {"Name": "validation_loss",
                                                       "Regex": ".*loss.*",
                                                       {"Name": "accuracy",
                                                          "Regex": ".*accuracy.*"}],
                             train_instance_count=1,
                             train_instance_type='ml.c4.xlarge')
```

## 2) Map S3 prefix to local download directory

The key name e.g. 'train' matches environment variable SM\_CHANNEL\_TRAIN suffix TRAIN, for train data directory

```
abalone_estimator.fit({'train': s3_input_prefix,
                       'validation': s3_input_prefix,
                       job_name="abalone-epoch-{}-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime()), epochs)})
```

# PyTorch training with SageMaker

## 3) Pass hyper parameters

The hyperparameter dict key name e.g. “traindata”, “epochs” matches the argument name “--traindata”, “--epochs”

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the")
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts',
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input")
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifacts',
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "."))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "."))

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker")

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the model",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.pytorch import PyTorch
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker.Session().bucket)
abalone_estimator = PyTorch(entry_point='main_train.py',
                             source_dir="./source",
                             role=role,
                             py_version="py3",
                             framework_version="1.0.0",
                             hyperparameters={'traindata': 'abalone_train.csv',
                                                'validationdata': 'abalone_test.csv',
                                                'epochs': 10,
                                                'batch-size': 32},
                             model_dir=s3_model_path,
                             metric_definitions=[{"Name": "mean_squared_error",
                                                    "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_error",
                                                    "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_percentage_error",
                                                    "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                             train_instance_count=1,
                             train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                       'validation': s3_input_prefix},
                      job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```



# PyTorch training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input",
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', "."))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "."))

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker",
                    default=os.environ.get('SM_MODEL_DIR', None))

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save snapshots",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.pytorch import PyTorch
from time import gmtime, strftime
s3_model_path = "s3://my-bucket/model_dir"
abalone_estimator =
```

4) Save artifacts to output  
Save model to output to dir pointed by environment variable SM\_MODEL\_DIR. Artifacts placed here are automatically uploaded to S3 and available during inference

```
epochs=10,
batch_size=32,
model_dir=s3_model_path,
metric_definitions=[{"Name": "mean_squared_error",
                      "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                     {"Name": "mean_absolute_error",
                      "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                     {"Name": "mean_absolute_percentage_error",
                      "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"},
                     ],
train_instance_count=1,
train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# PyTorch training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the training data")
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts',
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input file")
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifacts',
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "."))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "."))

parser.add_argument("--model_dir", help="Do not use this.. required for SageMaker")

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the model",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=100)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.pytorch import PyTorch
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = PyTorch(entry_point='main_train.py',
                             source_dir="./source",
                             role=role,
                             py_version="py3",
                             framework_version="1.0.0",
                             hyperparameters={'traindata': 'traindata',
                                                'validationdata': 'validationdata',
                                                'epochs': 10,
                                                'batch-size': 32},
                             model_dir=s3_model_path,
                             metric_definitions=[{"Name": "mean_squared_error",
                                                  "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_error",
                                                  "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                  {"Name": "mean_absolute_percentage_error",
                                                  "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                             train_instance_count=1,
                             train_instance_type='ml.c4.xlarge')
```

5) Specify the instance type for your training. Increase the instance count if your code supports distributed training

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="ablone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# General SageMaker Inference Flow

## Sample http request

POST /endpoints/abalone.. HTTP/1.1

Host: *runtime.sagemaker* ..

..

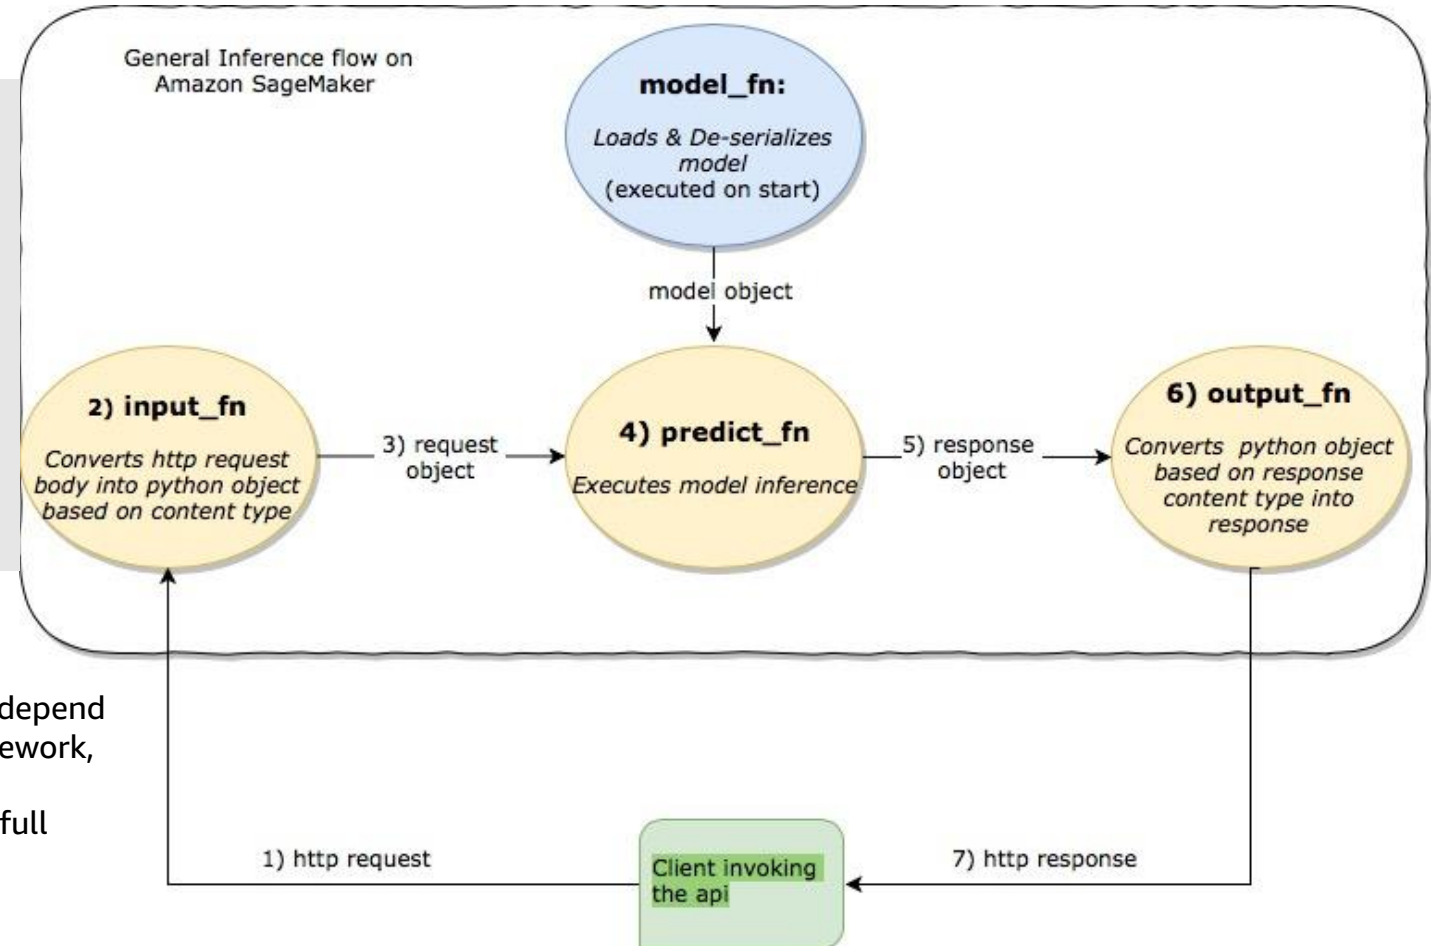
Authorization: AWS4-HMAC-SHA256 Credential \*\*\*

Content-Type: application/json

Accept: application/json

[[.34,5.6],[4,56]...]

**Note:** This is a general flow only. The exact function signatures depend on the SageMaker container for the specific deep learning framework, including its version! Please check the SageMaker samples on <https://github.com/aws-labs/amazon-sagemaker-examples> for full details



# PyTorch inference with SageMaker – Load model

```
def model_fn(model_dir):  
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
    model = torch.nn.DataParallel(Net())  
    with open(os.path.join(model_dir, 'model.pth'), 'rb') as f:  
        model.load_state_dict(torch.load(f))  
    return model.to(device)
```

Note: Your model can be deployed on CPU or GPU instance type you choose when you deploy the endpoint. This code supports both.

# Deploy my estimator to a SageMaker Endpoint and get a Predictor predictor  
= # E.g. m4.xlarge is a CPU instance type or if you use ml.p3.2xlarge then it is a GPU instance type

```
abalone_estimator.deploy(  
    instance_type='ml.m4.xlarge',  
    initial_instance_count=1)
```

# MXNet on SageMaker

# MXNet training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the training directory", required=True)
# The environment variable SM_CHANNEL_TRAIN is defined
parser.add_argument('--traindata-dir',
                    help='The directory containing training artifacts such as training data',
                    default=os.environ.get('SM_CHANNEL_TRAIN', "."))

# val dir files
parser.add_argument("--validationdata", help="The validation input file wrt to the val directory", required=True)
parser.add_argument('--validationdata-dir',
                    help='The directory containing validation artifacts such as validation data',
                    default=os.environ.get('SM_CHANNEL_VALIDATION', "."))

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', "result_data")
                    )

parser.add_argument("--model_dir", help="Do not use this.. required by SageMaker", default=None)

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the snapshot to..",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=10, type=int)
parser.add_argument("--batch-size", help="The mini batch size", default=30, type=int)
```



# MXNet training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', ".")

# val dir files
parser.add_argument("--validationdata", help="The validation input",
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', ".")

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', ".")
                    )

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker",
                    # This is where the model needs to be saved to
                    parser.add_argument("--snapshot_dir", help="The directory to save snapshots",
                                        default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=100)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.mxnet import MXNet
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = MXNet(entry_point='main_train.py',
                           source_dir="./source",
                           role=role,
                           py_version="py3",
                           framework_version="1.3.0",
                           hyperparameters={'traindata': 'abalone_train.csv',
                                             'validationdata': 'abalone_test.csv',
                                             'epochs': 10,
                                             'batch-size': 32},
                           model_dir=s3_model_path,
                           metric_definitions=[{"Name": "mean_squared_error",
                                                "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                              {"Name": "mean_absolute_error",
                                                "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                              {"Name": "mean_absolute_percentage_error",
                                                "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                           train_instance_count=1,
                           train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                       'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# MXNet training with SageMaker

1) Specify source code  
The entry point file and source code dir

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', ".")

# val dir files
parser.add_argument("--validationdata", help="The validation input",
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', ".")

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', ".")
                    )

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker",
                    default=os.environ.get('SM_MODEL_DIR', None))

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save snapshots",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=100)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.mxnet import MXNet
from time import gmtime, strftime
s3_model_path = "s3://[BUCKET]/models" format(sagemaker_session.default_bucket())
abalone_estimator = MXNet(entry_point='main_train.py',
                           source_dir='./source',
                           role=role,
                           py_version="py3",
                           framework_version="1.3.0",
                           hyperparameters={'traindata': 'abalone_train.csv',
                                              'validationdata': 'abalone_test.csv',
                                              'epochs': 10,
                                              'batch-size': 32},
                           model_dir=s3_model_path,
                           metric_definitions=[{"Name": "mean_squared_error",
                                                  "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                {"Name": "mean_absolute_error",
                                                  "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                {"Name": "mean_absolute_percentage_error",
                                                  "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                           train_instance_count=1,
                           train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                       'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# MXNet training with SageMaker

```
# Train dir files
parser.add_argument("--traindata", help="The input file wrt to the",
                    # The environment variable SM_CHANNEL_TRAIN is defined
                    parser.add_argument('--traindata-dir',
                                        help='The directory containing training artifacts',
                                        default=os.environ.get('SM_CHANNEL_TRAIN', ".")

# val dir files
parser.add_argument("--validationdata", help="The validation input",
                    parser.add_argument('--validationdata-dir',
                                        help='The directory containing validation artifacts',
                                        default=os.environ.get('SM_CHANNEL_VALIDATION', ".")

# output dir to save any files such as predictions, logs, etc
parser.add_argument("--outputdir", help="The output dir to save results",
                    default=os.environ.get('SM_OUTPUT_DATA_DIR', ".")
                    )

parser.add_argument("--model_dir", help="Do not use this.. requires SageMaker",
                    default=os.environ.get('SM_MODEL_DIR', None))

# This is where the model needs to be saved to
parser.add_argument("--snapshot_dir", help="The directory to save the model",
                    default=os.environ.get('SM_MODEL_DIR', None))

# Additional parameters for your code
parser.add_argument("--epochs", help="The number of epochs", default=100)
parser.add_argument("--batch-size", help="The mini batch size", default=32)
```

## Submit your training job

```
from sagemaker.mxnet import MXNet
from time import gmtime, strftime
s3_model_path = "s3://{}/models".format(sagemaker_session.default_bucket())
abalone_estimator = MXNet(entry_point='main_train.py',
                           source_dir="./source",
                           role=role,
                           py_version="py3",
                           framework_version="1.3.0",
                           hyperparameters={'traindata': 'abalone_train.csv',
                                             'validationdata': 'abalone_test.csv',
                                             'epochs': 10,
                                             'batch-size': 32},
                           model_dir=s3_model_path,
                           metric_definitions=[{"Name": "mean_squared_error",
                                                "Regex": "## validation_metric_mse ##: (\\d*[.]?\\d*)"},
                                                {"Name": "mean_absolute_error",
                                                "Regex": "## validation_metric_mae ##: (\\d*[.]?\\d*)"},
                                                {"Name": "mean_absolute_percentage_error",
                                                "Regex": "## validation_metric_mape ##: (\\d*[.]?\\d*)"}],
                           train_instance_count=1,
                           train_instance_type='ml.c4.xlarge')
```

```
abalone_estimator.fit({'train': s3_input_prefix,
                      'validation': s3_input_prefix},
                      job_name="abalone-age-py3-{}".format(strftime("%Y-%m-%d-%H-%M-%S", gmtime())))
```

# MXNet inference with SageMaker

```
def model_fn(model_dir):  
    """  
    Load the gluon model. Called once when hosting service starts.  
    :param: model_dir The directory where model files are stored.  
    :return: a model (in this case a Gluon network)  
    """  
    net = gluon.SymbolBlock.imports(  
        '%s/model-symbol.json' % model_dir,  
        ['data'],  
        '%s/model-0000.params' % model_dir,  
    )  
    return net
```

# Deploy your MXNet estimator to a SageMaker Endpoint

```
abalone_estimator.deploy(  
    instance_type='ml.m4.xlarge',  
    initial_instance_count=1)
```

1. No code required if using default SageMaker MXNet Model Server.
2. Works well if the model is a single artifact of type MXNet nn.module
3. Else write custom code, e.g load Gluon model..

# MXNet inference with SageMaker – Custom Dataformats

```
def input_fn(request_body, request_content_type, model):  
    """An input_fn that loads a pickled numpy array"""  
    if request_content_type == "application/python-pickle":  
        array = np.load(StringIO(request_body))  
        array.reshape(model.data_shapes[0])  
        return mx.io.NDArrayIter(mx.ndarray(array))  
    else:  
        # Handle other content-types here or raise an Exception  
        # if the content type is not supported.  
        pass
```

**Note:** This sample is for SageMaker with MXNet 1.3.0

# Deserialize the Invoke request body into an object we can perform prediction on  
input\_object = input\_fn(request\_body, request\_content\_type, model)

# Perform prediction on the deserialized object, with the loaded model  
prediction = predict\_fn(input\_object, model)

# Serialize the prediction result into the desired response content type  
output = output\_fn(prediction, response\_content\_type)

# Summary – Amazon SageMaker

- Built-in support for opensource ML Frameworks TensorFlow, PyTorch, Chainer, MXNet and Scikit-learn
- Simplifies porting existing code to use AWS Platform, making your code fully portable
- Simplifies model endpoint inference creation – One Click Deployment
- Automatic monitoring and tracking logs, infrastructure and model performance
- Automatic experiment tracking including data, code, hyperparameters & logs



# Learn why customers choose AWS for machine learning



## Demo Arena

Watch how machine learning is used in real life applications



## Ask the Experts

Get your questions answered by AWS Experts



## Machine Learning on AWS

<https://aws.amazon.com/machine-learning/>

# Learn from AWS experts. Advance your skills and knowledge. Build your future in the AWS Cloud.



## Digital Training

Free, self-paced online courses built by AWS experts



## Classroom Training

Classes taught by accredited AWS instructors



## AWS Certification

Exams to validate expertise with an industry-recognized credential

Ready to begin building your cloud skills?  
Get started at: <https://www.aws.training/>

# Thank You for Attending AWS Innovate

We hope you found it interesting! A kind reminder to **complete the survey**.  
Let us know what you thought of today's event and how we can improve the event experience for you in the future.



[aws-apac-marketing@amazon.com](mailto:aws-apac-marketing@amazon.com)



[twitter.com/AWSCloud](https://twitter.com/AWSCloud)



[facebook.com/AmazonWebServices](https://facebook.com/AmazonWebServices)



[youtube.com/user/AmazonWebServices](https://youtube.com/user/AmazonWebServices)



[slideshare.net/AmazonWebServices](https://slideshare.net/AmazonWebServices)



[twitch.tv/aws](https://twitch.tv/aws)